

**Akademia Górniczo-Hutnicza
im. Stanisława Staszica w Krakowie**

Wydział Elektrotechniki, Automatyki, Informatyki i Elektroniki

KATEDRA INFORMATYKI STOSOWANEJ



PRACA MAGISTERSKA

KONRAD MALAWSKI

**PRZETWARZANIE I ANALIZA DANYCH
MULTIMEDIALNYCH W ŚRODOWISKU
ROZPROSZONYM**

PROMOTOR:

dr inż. Sebastian Ernst

Kraków 2013

OŚWIADCZENIE AUTORA PRACY

OŚWIADCZAM, ŚWIADOMY ODPOWIEDZIALNOŚCI KARNEJ ZA
POŚWIADCZENIE NIEPRAWDY, ŻE NINIEJSZĄ PRACĘ DYPLOMOWĄ
WYKONAŁEM OSOBIŚCIE I SAMODZIELNIE, I NIE KORZYSTAŁEM ZE
ŹRÓDEŁ INNYCH NIŻ WYMIENIONE W PRACY.

.....

PODPIS

AGH
University of Science and Technology in Krakow

Faculty of Electrical Engineering, Automatics, Computer Science and
Electronics

DEPARTMENT OF APPLIED COMPUTER SCIENCE



MASTER OF SCIENCE THESIS

KONRAD MALAWSKI

**PROCESSING AND ANALISYS OF MULTIMEDIA IN
DISTRIBUTED SYSTEMS**

SUPERVISOR:

Sebastian Ernst Ph.D

Krakow 2013

Acknowledgements...

Contents

1. Problem definition	7
1.1. Use cases	7
2. Preparing the cluster	9
3. Architecture overview	10
3.1. Hadoop Distributed File System - HDFS	10
3.2. Distributed Actor System - Akka	10
4. System design	12
4.1. Loader	12
4.1.1. Work sharing in an actor-based cluster.....	12
4.2. Job Runner	13
4.2.1. Tackling the „small-files problem” on HDFS.....	13
5. Results and processing times	15
6. Conclusions	16

Todo list

need ref	10
TODO explain why HDFS is the reasonable choice here	10
image for actor system here	10
TODO why do we need actors here? why does it make sense to distribute the work . . .	11
finish me	13

1. Problem definition

The primary goal of this work is to research how to efficiently work with large amounts of multimedia data in distributed systems. In order to guarantee the devised approach would make sense

1.1. Use cases

One of the simplest use cases in which this system might be used is trivial *duplicate detection* of video files. One might think that a simple checksum of the video files would suffice, but imagine a system like *YouTube* where hours of video content are uploaded each second. It's easy to imagine two people uploading the same trailer for an upcoming movie during the same day. In the easiest case, a simple checksum check would suffice, but what if the videos were encoded using different codecs or even were uploaded in different resolutions?

To further represent how difficult it may be to answer the question of „Are those two files the same material?“ let's go back to the example of users uploading video content to our YouTube-like service. What if the users knew the uploaded material is copy righted and should not be uploaded to the public internet? As one might imagine, this might happen quite often on a service like YouTube. Let's assume our users know about our „duplicate“ detection algorithm and that we might be partnering with movie companies, so that we have reference material on the servers we can compare the uploaded videos to. Obviously the users will try to trick our algorithm into not recognising that the uploaded content is the same as our reference material. One of the tricks often seen on YouTube.com is that the users upload the material „mirrored“ which further complicates our matching algorithm – we must now also be resilient against data modified especially for the goal of us not being able to detect it.

Another, less copyright focused, goal of the presented system is to be able to extensively mine data from the uploaded video content. Here the canonical example would be a „*TOP 10 Movies of All Time*“ video, which obviously contains video material from at least 10 movies, usually in the order of 10th, 9th ... until the 1st (best) movie of all time. If we would be able to match parts of each video to their corresponding reference materials, we would be able to get meta data about the now recognised movies and even mine out the data what is the best / worst

movie of all time, even without it being written per se - only by looking at the frames in the video. This idea only scratches the surface of what the system implemented during this thesis work might do, but it will be our test case that we'll be working towards during this paper.

The system implemented during this work represents a basic effort to tackle the above problems as well as these goals

2. Preparing the cluster

Hadoop's filesystem must be formatted before put into use. This is achieved by issuing the `-format` command to the namenode:

```
kmalawski@oculus-master > hadoop namenode -format
```

It is worth pointing out that a "format" takes place only on the namenode, it does not actually touch the data stored on the datanodes, but instead it deleted the data stored on the namenode. The Namenode, as explained previously, stores all metadata about where a file is located, thus, cleaning it's data makes the files stores in HDFS un-usable, since we don't know "where a file's chunks are stored".

3. Architecture overview

In this chapter I will present a broad overview on the system's design and various components. It should also provide some background to why a distributed file system was required in order to enable this system, and why only a distributed system is suited to handle the kind of jobs that Oculus is designed for.

3.1. Hadoop Distributed File System - HDFS

The primary enabler for this project, and many „big data” projects in the recent years is Hadoop and it's Distributed File System – for short *HDFS*.

Hadoop is a Java implementation of the „Map Reduce” white paper published in 2004 by Google . It first developed internally at *Yahoo* and then open-sourced in 2007 and is now under the Apache Software Foundation's umbrella,

need
ref

3.2. Distributed Actor System - Akka

In order to enable fast downloading and uploading of video material into the system, an Actor System based solution was implemented. The used framework, which provided the basic Actor building blocks, is Akka, a JVM based implementation of the Erlang Actor Model of Concurrency. The implementation language selected was Scala, as it's a first-class citizen in the Akka world, as well as for it's conciseness and semantic power.

TODO
ex-
plain
why

HDFS
is
the
rea-
son-
able
choice
here

The Actor Model of computation raises the level of abstraction we operate on in parallel applications beyond threads, and instead we use Actors. An Actor can be described as an entity that can only receive and send messages. An external system (the so-called "Actor System") is responsible for assigning Threads to Actors, so that they can perform the work - upon receiving a message.

image
for
actor
sys-
tem
here

This model has the advantage of hiding any state the Actor may have from other Actors it interacts with, thus there are by definition no race conditions as no shared-mutable state can

be observed in such a system. While the Actor Model of concurrency was designed for local interactions - that is, avoiding shared state in highly parallel systems it also has a significant impact on distributed systems. Because actors share no state, and it is only possible to interact with an Actor using messages it does not matter *where* an Actor is residing – for example it might reside on another node in our computation cluster. This property was used in Oclous in order to distribute the work-load related to downloading, pre-processing as well as initially uploading the raw video content into the distributed file system.

TODO

why
do
we
need
ac-
tors
here?
why
does
it
make
sense
to
dis-
tribute
the
work

4. System design

The system, from here to be referred to by the name „*Oculus*”, is designed with an asynchronous as well as distributed approach in mind. In order to achieve high asynchronicity between obtaining new reference data, and running jobs such as „compare video1 with the reference database”, the system was split into two primary components:

- **loader** – which is responsible for obtaining more and more reference material. It persists and initially processes the videos, as well as any related metadata,
- and the **job runner** – which is responsible for running computations on top of the hadoop cluster and reference databases.

In this chapter I will discuss the high-level design of these subsystems, and how they interact with each other. In the next chapter I will describe each of the systems technical challenges and how they were solved.

4.1. Loader

The downloader is responsible for obtaining as much as possible „reference data”, by which I mean video material – from sites such as *youtube.com* or video hosting sites.

It should be also noted that while I refer to this system as „the” downloader, in fact it is composed out of many instances of the same application, sharing the common workload of downloading and converting the input data. The deployment diagram on drawing [12] explains the interactions between the nodes.

4.1.1. Work sharing in an actor-based cluster

The downloader system is implemented in an „message passing” style, which means that the work requests to the cluster come in as *messages*, and then are *routed* to an *actor*, who performs the work, and then responds with another message – in other words, all communication between Actor instances is performed via messages and there is no shared state between them. This also applies to Actor instances residing in the same JVM - for the user of an „Actor System”,

the location (on which physical node the actor is actually executing) of an actor remains fully transparent - this has huge gains in terms of load balancing the message execution on the cluster - the router decides who will be notified on which message, the listing 4.1 shows a typical workflow using the so-called „smallest inbox” routing strategy ??.

Listing 4.1: smallest-inbox routing algorithm

```

                                ---> [inbox1, size = 3]
                                =>
                                YoutubeDownloadActor(1)
YoutubeCrawlerActor -- Msg(url=http://...) --> router ---/
                                \   [inbox2, size = 5]
                                =>
                                YoutubeDownloadActor(2)

                                /* because inbox1.size < inbox2.size */

```

The underlying Actor System is implemented by a project called Akka ??, and can be easiest explained as „porting Erlang concepts of the Actor Model to the JVM”. I selected the „Smallest Inbox” routing strategy instead of the other widely used „Round Robin” approach in order to guarantee not overloading any Actor with too many requests to download movies (which is a relatively slow process). Thanks to the smallest inbox routing, I can guarantee that if some of the nodes have a faster connection to the Internet, they will get more movie download requests, than nodes located on a slower network.

As mentioned before, the system is fully distributed and *any node can perform any task* submitted to the cluster. For example let’s take the first step in the processing pipeline in Oculus, which is Download video from <http://www.youtube.com/watch?v=-X9bcrJ3TjY> – such message will be emitted by the YoutubeCrawlerActor and sent via a Router instance to an YoutubeDownloadActor which has the „smallest inbox”.

4.2. Job Runner

The job runners responsibility is at the core of the systems...

finish
me

4.2.1. Tackling the „small-files problem” on HDFS

In this section I will explain the „small-files problem” which would lead to major performance degradation of the Hadoop cluster is left undressed, and how I solved it in the reference system – Oculus.

The so called „*small files problem*” is a known problem in the Hadoop world relating to how HDFS stores files. The problem arises whenever we try to store files smaller than the block size used by HDFS to allocate files. One can think of it as wasting space – because Hadoop allocates at least one block per file we store in it, and all operations (such as migrating data between nodes, reading files into memory etc.) are optimised towards larger block-sizes (usually around 32MB or 64MB).

To illustrate the problem with an example: if we try storing a 2MB file in HDFS we’ll in fact produce 62MB of „wasted space”, reducing the cluster’s capacity way faster than we ought to. Another p

5. Results and processing times

6. Conclusions

The applied technologies have indeed been very helpful, and proved to be very elastic for different kinds of jobs related to processing large amounts of data. I was also positively surprised with the ease of Scaling Hadoop infrastructure.

Bibliography